

In [232]:

```
# Spring 2022
# IE7275- Data Mining in Engineering
# Project Group 24
# Authors - Samruddhi Kulkarni
#           Karthik Vadlamani

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn import preprocessing
from sklearn.model_selection import train_test_split
```

In [233]:

```
# Importing the csv file of the dataset
stroke_df = pd.read_csv("C:/Sam/Sem2/Data mining/Project/Healthcare dataset/archive/healthcare-dataset-stroke-data.csv")
print(stroke_df)
```

	id	gender	age	hypertension	heart_disease	ever_married	\
0	9046	Male	67.0	0	1	Yes	
1	51676	Female	61.0	0	0	Yes	
2	31112	Male	80.0	0	1	Yes	
3	60182	Female	49.0	0	0	Yes	
4	1665	Female	79.0	1	0	Yes	
...	...	...	...	...	...	...	
5105	18234	Female	80.0	1	0	Yes	
5106	44873	Female	81.0	0	0	Yes	
5107	19723	Female	35.0	0	0	Yes	
5108	37544	Male	51.0	0	0	Yes	
5109	44679	Female	44.0	0	0	Yes	

	work_type	Residence_type	avg_glucose_level	bmi	smoking_status	\
0	Private	Urban	228.69	36.6	formerly smoked	
1	Self-employed	Rural	202.21	NaN	never smoked	
2	Private	Rural	105.92	32.5	never smoked	
3	Private	Urban	171.23	34.4	smokes	
4	Self-employed	Rural	174.12	24.0	never smoked	
...	...	...	...	...	...	
5105	Private	Urban	83.75	NaN	never smoked	
5106	Self-employed	Urban	125.20	40.0	never smoked	
5107	Self-employed	Rural	82.99	30.6	never smoked	
5108	Private	Rural	166.29	25.6	formerly smoked	
5109	Govt_job	Urban	85.28	26.2	Unknown	

	stroke
0	1
1	1
2	1
3	1
4	1
...	...
5105	0
5106	0
5107	0
5108	0
5109	0

[5110 rows x 12 columns]

In [234]:

```
# Printing the descriptive statistics of the dataset
stroke_df.describe()
```

Out [234]:

	id	age	hypertension	heart_disease	avg_glucose_level	bmi	stroke
count	5110.000000	5110.000000	5110.000000	5110.000000	5110.000000	4909.000000	5110.000000
mean	36517.829354	43.226614	0.097456	0.054012	106.147677	28.893237	0.048728
std	21161.721625	22.612647	0.296607	0.226063	45.283560	7.854067	0.215320
min	67.000000	0.080000	0.000000	0.000000	55.120000	10.300000	0.000000
25%	17741.250000	25.000000	0.000000	0.000000	77.245000	23.500000	0.000000
50%	36932.000000	45.000000	0.000000	0.000000	91.885000	28.100000	0.000000
75%	54682.000000	61.000000	0.000000	0.000000	114.090000	33.100000	0.000000
max	72940.000000	82.000000	1.000000	1.000000	271.740000	97.600000	1.000000

In [235]:

```
# Printing the information of column name and datatype
stroke_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5110 entries, 0 to 5109
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                     5110 non-null   int64
1   gender                 5110 non-null   object
2   age                    5110 non-null   float64
3   hypertension           5110 non-null   int64
4   heart_disease          5110 non-null   int64
5   ever_married           5110 non-null   object
6   work_type              5110 non-null   object
7   Residence_type         5110 non-null   object
8   avg_glucose_level      5110 non-null   float64
9   bmi                    4909 non-null   float64
10  smoking_status         5110 non-null   object
11  stroke                 5110 non-null   int64
dtypes: float64(3), int64(4), object(5)
memory usage: 479.2+ KB
```

In [236]:

```
# Checking for null values in the dataset
stroke_df.isna().sum()
```

Out[236]:

```
id                0
gender            0
age              0
hypertension      0
heart_disease     0
ever_married      0
work_type         0
Residence_type    0
avg_glucose_level 0
bmi              201
smoking_status    0
stroke            0
dtype: int64
```

In [237]:

```
# We see that bmi has 201 NA rows, which we can fill with median values
stroke_df['bmi'].fillna(stroke_df['bmi'].median(), inplace=True)
# Checking again for null values
stroke_df.isna().sum()
# The output confirms all the null values were filled with median values
```

Out[237]:

```
id 0
gender 0
age 0
hypertension 0
heart_disease 0
ever_married 0
work_type 0
Residence_type 0
avg_glucose_level 0
bmi 0
smoking_status 0
stroke 0
dtype: int64
```

In [238]:

```
# Identifying and removing redundant data
# We can drop the id column as it is redundant
stroke_df.drop('id',axis=1,inplace=True)
stroke_df
```

Out[238]:

	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose_level	bmi	smoking_status
0	Male	67.0	0	1	Yes	Private	Urban	228.69	36.6	formerly smoked
1	Female	61.0	0	0	Yes	Self-employed	Rural	202.21	28.1	never smoked
2	Male	80.0	0	1	Yes	Private	Rural	105.92	32.5	never smoked
3	Female	49.0	0	0	Yes	Private	Urban	171.23	34.4	formerly smoked
4	Female	79.0	1	0	Yes	Self-employed	Rural	174.12	24.0	never smoked
...	...	...	...	...	...	...	...	...	...	...
5105	Female	80.0	1	0	Yes	Private	Urban	83.75	28.1	never smoked
5106	Female	81.0	0	0	Yes	Self-employed	Urban	125.20	40.0	never smoked
5107	Female	35.0	0	0	Yes	Self-employed	Rural	82.99	30.6	never smoked
5108	Male	51.0	0	0	Yes	Private	Rural	166.29	25.6	formerly smoked
5109	Female	44.0	0	0	Yes	Govt_job	Urban	85.28	26.2	Unknown

5110 rows x 11 columns



In [239]:

```
# Determining unique values of all categorical columns
print(stroke_df.gender.unique())
print(stroke_df.ever_married.unique())
print(stroke_df.work_type.unique())
print(stroke_df.Residence_type.unique())
print(stroke_df.smoking_status.unique())
```

```
['Male' 'Female' 'Other']
['Yes' 'No']
['Private' 'Self-employed' 'Govt_job' 'children' 'Never_worked']
['Urban' 'Rural']
['formerly smoked' 'never smoked' 'smokes' 'Unknown']
```

In [240]:

```
# From the above categories we can see that the 'Other' category in Gender field has just
```

```

one entry which we can remove
stroke_df.drop(stroke_df.index[stroke_df['gender'] == 'Other'], inplace = True)
print(stroke_df.gender.unique())

```

```
['Male' 'Female']
```

In [241]:

```

# Finding duplicate rows in the dataset
duplicate = stroke_df[stroke_df.duplicated()]

print("Duplicate Rows :")
duplicate
stroke_df
# From the output it is clear that there are no duplicate data in the dataset

```

Duplicate Rows :

Out[241]:

	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose_level	bmi	smoking_status
0	Male	67.0	0	1	Yes	Private	Urban	228.69	36.6	for sn
1	Female	61.0	0	0	Yes	Self-employed	Rural	202.21	28.1	never sn
2	Male	80.0	0	1	Yes	Private	Rural	105.92	32.5	never sn
3	Female	49.0	0	0	Yes	Private	Urban	171.23	34.4	sr
4	Female	79.0	1	0	Yes	Self-employed	Rural	174.12	24.0	never sn
...	...	...	...	...	...	...	...	...	...	
5105	Female	80.0	1	0	Yes	Private	Urban	83.75	28.1	never sn
5106	Female	81.0	0	0	Yes	Self-employed	Urban	125.20	40.0	never sn
5107	Female	35.0	0	0	Yes	Self-employed	Rural	82.99	30.6	never sn
5108	Male	51.0	0	0	Yes	Private	Rural	166.29	25.6	for sn
5109	Female	44.0	0	0	Yes	Govt_job	Urban	85.28	26.2	Unk

5109 rows x 11 columns

In [242]:

```


# Plotting barplots of stroke vs categorical fields
categorical = stroke_df.select_dtypes('object').columns
print(categorical)
plt.figure(figsize = (20,20))
i = 1
for column in categorical[:-1]:
    plt.subplot(4, 2, i)
    sns.countplot(x = stroke_df[column], hue = stroke_df["stroke"])
    i+=1
plt.show()

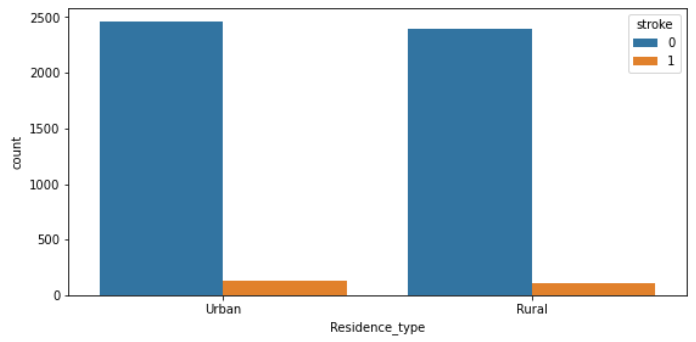
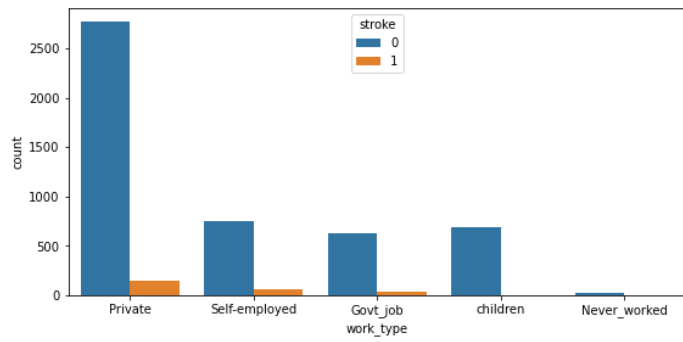


```

```

Index(['gender', 'ever_married', 'work_type', 'Residence_type',
      'smoking_status'],
      dtype='object')


```






In [243]:


```
# Plotting barplot of stroke vs smoking_status field
sns.countplot(x = stroke_df["smoking_status"], hue = stroke_df["stroke"])
plt.show()
```



In [244]:

```
# Plotting boxplots of numerical fields like age, bmi, average glucose level of patients
plt.figure(figsize = (10,8))
df2 = pd.DataFrame(data=stroke_df, columns=["age", "bmi", "avg_glucose_level"])
sns.boxplot(x="variable", y="value", data=pd.melt(df2))
plt.show()
```





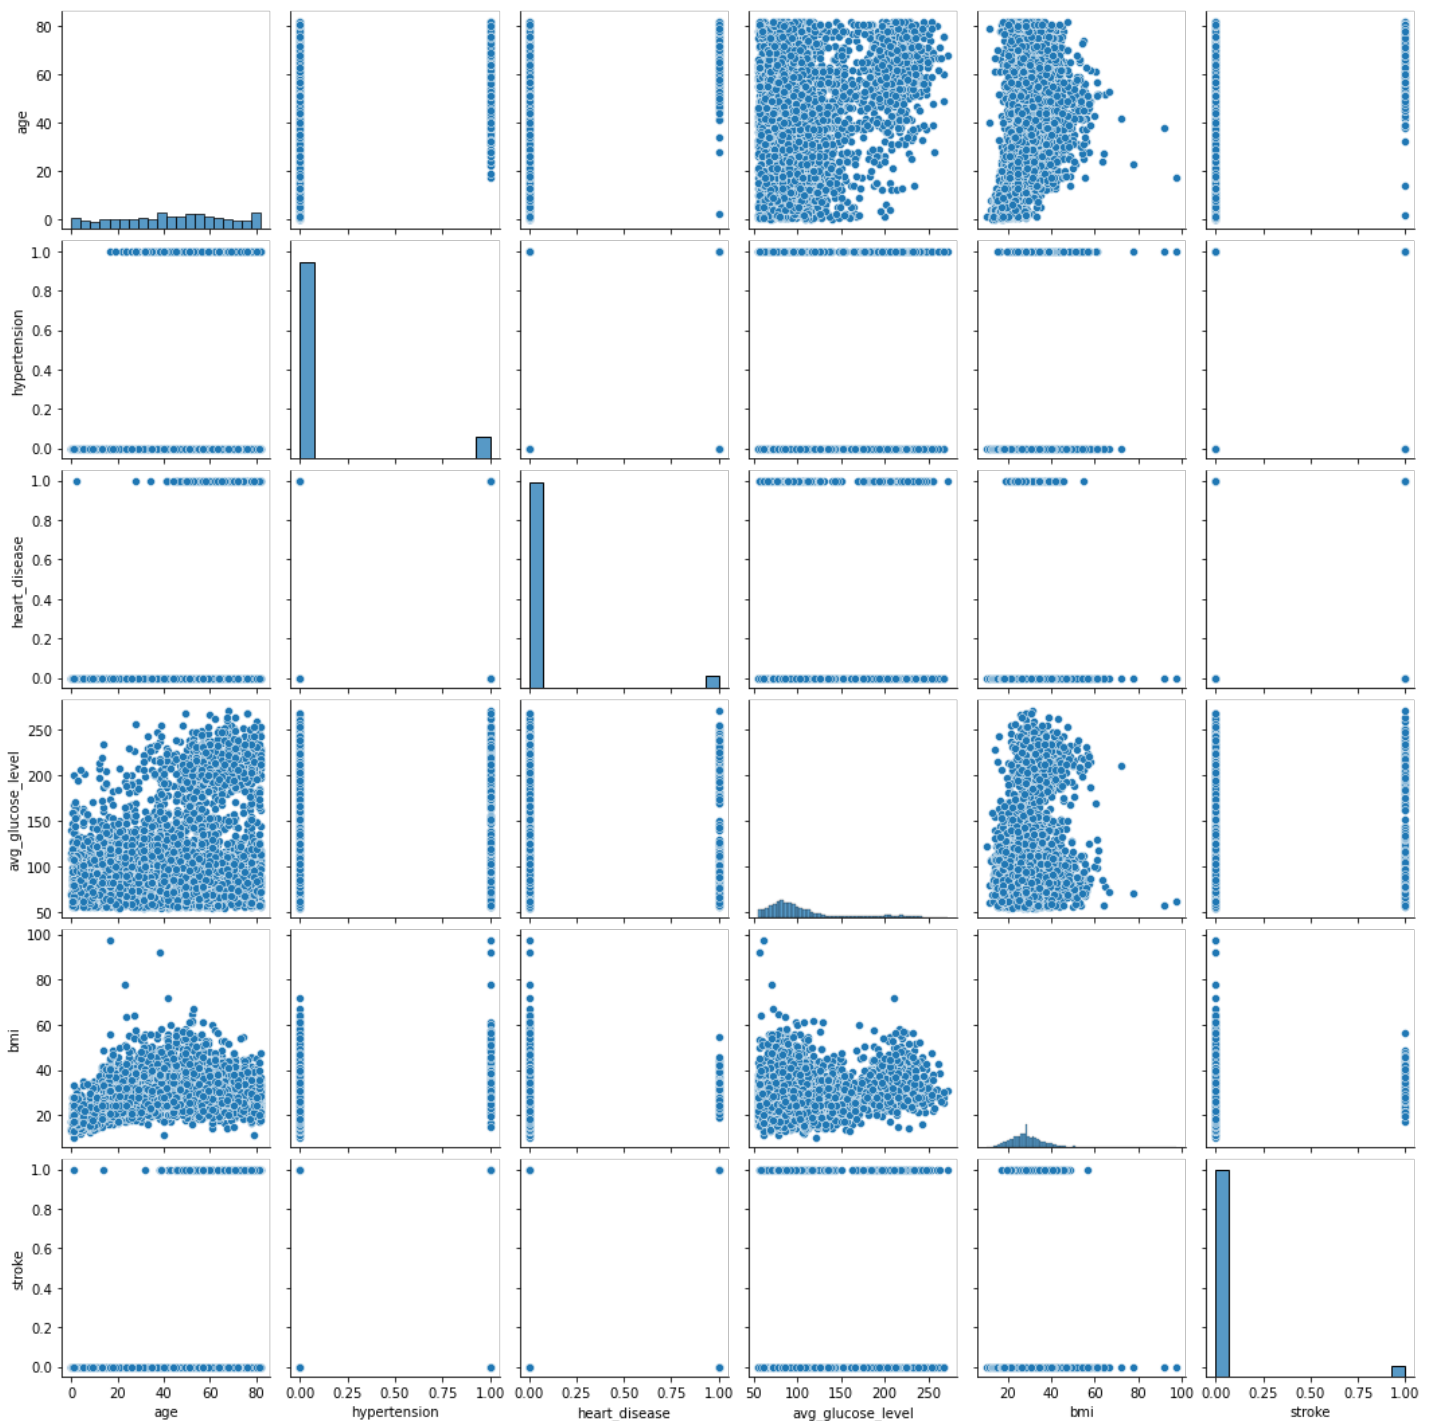
In [245]:

```
# Plotting pairplots the fields
plt.figure(figsize = (15,15))
sns.pairplot(stroke_df)
```

Out[245]:

<seaborn.axisgrid.PairGrid at 0x1a58439fbe0>

<Figure size 1080x1080 with 0 Axes>



In [246]:

```
!pip install plotly
```

Requirement already satisfied: plotly in c:\users\samru\anaconda3\lib\site-packages (5.7.0)

Requirement already satisfied: six in c:\users\samru\anaconda3\lib\site-packages (from plotly) (1.16.0)

Requirement already satisfied: tenacity>=6.2.0 in c:\users\samru\anaconda3\lib\site-packages (from plotly) (6.2.0)

Requirement already satisfied: tenacity<=0.2.0 in c:\users\samir\anaconda3\lib\site-packages (from plotly) (8.0.1)

In [247]:

```
# Plotting scatterplot of age vs stroke
import plotly.express as px
plt.figure(figsize = (5,5))
fig = px.scatter( stroke_df, x = 'age', y = 'stroke')
fig.show()
```

<Figure size 360x360 with 0 Axes>

In [248]:

```
# Some fields in the dataset have object datatype
# For applying k-NN to categorical variables, we need to convert them to binary dummy variables
stroke_df.gender = stroke_df.gender.astype('category').cat.codes
stroke_df.ever_married = stroke_df.ever_married.astype('category').cat.codes
stroke_df.Residence_type = stroke_df.Residence_type.astype('category').cat.codes
stroke_df = pd.get_dummies(data=stroke_df, prefix=['work_type', 'smoking_status'], columns=['work_type', 'smoking_status'], drop_first=True)
stroke_df.head(10)
#print(stroke_df[['age', 'bmi', 'avg_glucose_level']])
```

Out[248]:

	gender	age	hypertension	heart_disease	ever_married	Residence_type	avg_glucose_level	bmi	stroke	work_type_Never
0	1	67.0	0	1	1	1	228.69	36.6	1	
1	0	61.0	0	0	1	0	202.21	28.1	1	
2	1	80.0	0	1	1	0	105.92	32.5	1	
3	0	49.0	0	0	1	1	171.23	34.4	1	
4	0	79.0	1	0	1	0	174.12	24.0	1	
5	1	81.0	0	0	1	1	186.21	29.0	1	

6	gender	1	74.0	hypertension	1	heart_disease	1	ever_married	1	Residence_type	0	avg_glucose_level	70.09	bmi	27.4	stroke	1	work_type_Never
7	0	69.0		0		0		0		1		94.39	22.8			1		
8	0	59.0		0		0		1		0		76.15	28.1			1		
9	0	78.0		0		0		1		1		58.57	24.2			1		

In [249]:

```
print(stroke_df[['age', 'bmi', 'avg_glucose_level']].describe())
```

```

count    5109.000000    5109.000000    5109.000000
mean      43.229986     28.863300     106.140399
std       22.613575      7.699785      45.285004
min        0.080000     10.300000     55.120000
25%       25.000000     23.800000     77.240000
50%       45.000000     28.100000     91.880000
75%       61.000000     32.800000    114.090000
max       82.000000     97.600000    271.740000

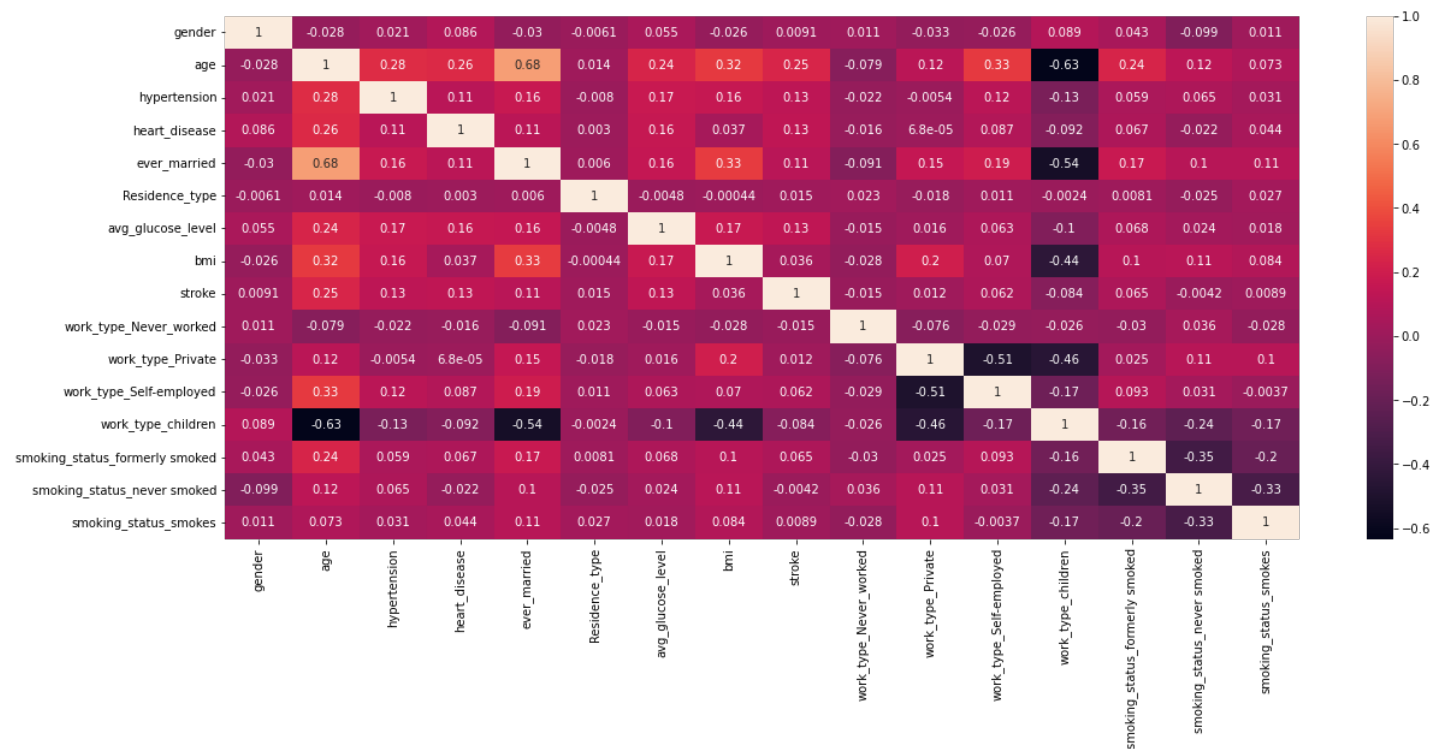
```

In [278]:

```

# Plotting heatmap of correlation of fields in the dataset
corrMatrix = stroke_df.corr()
# print (corrMatrix)
plt.figure(figsize=(20,8))
sns.heatmap(corrMatrix, annot=True)
plt.show()

```



In [251]:


```

# Pieplot of Hypertension patients distribution
fig, axes = plt.subplots(figsize=(8,8))
stroke_df['hypertension'].value_counts().plot.pie(autopct='%1.1f%%',
                                                    colormap='Pastel2',
                                                    fontsize=20,
                                                    shadow=True,
                                                    explode=[0.2,0])
axes.set_title("Pie Chart Of Hypertension", fontsize=20)
axes.set_xlabel(None)
plt.show()

```

Pie Chart Of Hypertension






In [252]:

```
# Pieplot of Heart disease patients distribution
fig, axes = plt.subplots(figsize=(8,8))
stroke_df['heart_disease'].value_counts().plot.pie(autopct='%1.1f%%',
                                                    colormap='Pastel2',
                                                    fontsize=20,
                                                    shadow=True,
                                                    explode=[0.2,0])

axes.set_title("Pie Chart Of Heart Disease", fontsize=20)
axes.set_xlabel(None)
plt.show()
```

Pie Chart Of Heart Disease



In [253]:

```
#Standardizing the numerical columns
```

```
col = ['age', 'avg_glucose_level', 'bmi']
min_max_scaler = preprocessing.MinMaxScaler()

column_names_to_normalize = ['age', 'avg_glucose_level', 'bmi']
x = stroke_df[column_names_to_normalize].values
x_scaled = min_max_scaler.fit_transform(x)
df_temp = pd.DataFrame(x_scaled, columns=column_names_to_normalize, index = stroke_df.index)
stroke_df[column_names_to_normalize] = df_temp
# Verifying the operation through describe function
stroke_df.describe()
print(stroke_df[['age', 'bmi', 'avg_glucose_level']].describe())
```

```
count    age    bmi    avg_glucose_level
count  5109.000000  5109.000000  5109.000000
mean    0.526733    0.212638    0.235529
std     0.276045    0.088199    0.209053
min     0.000000    0.000000    0.000000
25%    0.304199    0.154639    0.102114
50%    0.548340    0.203895    0.169698
75%    0.743652    0.257732    0.272228
max     1.000000    1.000000    1.000000
```

In [254]:

```
# Forming feature and target dataframes
X = stroke_df.drop(['stroke'], axis=1)
y = stroke_df['stroke']
```

In [255]:

```
# Splitting data as training = 75% and testing = 25%
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42)
```

In [256]:

```
# transform the dataset by oversampling
from imblearn.over_sampling import SMOTE
oversample = SMOTE()
X_oversample, y_oversample = oversample.fit_resample(X, y)
X_train_os, X_test_os, y_train_os, y_test_os = train_test_split(X_oversample, y_oversample, test_size=0.20, random_state=42)
```

In [257]:

```
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from sklearn.neighbors import NearestNeighbors, KNeighborsClassifier

# KNN Classifier before oversampling
results = []
for k in range(1, 40):
    knn = KNeighborsClassifier(n_neighbors=k).fit(X_train, y_train)
    results.append({'k': k, 'accuracy': accuracy_score(y_test, knn.predict(X_test)) })
# Convert results to a pandas data frame
results = pd.DataFrame(results)
print(results)

# KNN Classifier after oversampling
results1 = []
print('After oversampling')
for k in range(1, 40):
    knn = KNeighborsClassifier(n_neighbors=k).fit(X_train_os, y_train_os)
    results1.append({'k': k, 'accuracy': accuracy_score(y_test_os, knn.predict(X_test_os)) })
# Convert results to a pandas data frame
results1 = pd.DataFrame(results1)
print(results1)
```

```
   k  accuracy
0   1  0.904538
1   2  0.925055
```

1	2	0.935055
2	3	0.928013
3	4	0.938185
4	5	0.937402
5	6	0.938185
6	7	0.937402
7	8	0.938185
8	9	0.938967
9	10	0.938185
10	11	0.938967
11	12	0.937402
12	13	0.938185
13	14	0.937402
14	15	0.937402
15	16	0.937402
16	17	0.937402
17	18	0.937402
18	19	0.937402
19	20	0.937402
20	21	0.937402
21	22	0.937402
22	23	0.937402
23	24	0.937402
24	25	0.937402
25	26	0.937402
26	27	0.937402
27	28	0.937402
28	29	0.937402
29	30	0.937402
30	31	0.937402
31	32	0.937402
32	33	0.937402
33	34	0.937402
34	35	0.937402
35	36	0.937402
36	37	0.937402
37	38	0.937402
38	39	0.937402

After oversampling

	k	accuracy
0	1	0.899691
1	2	0.889918
2	3	0.887346
3	4	0.889403
4	5	0.873971
5	6	0.874486
6	7	0.867284
7	8	0.871399
8	9	0.860597
9	10	0.866255
10	11	0.858539
11	12	0.858025
12	13	0.850309
13	14	0.852366
14	15	0.843621
15	16	0.844136
16	17	0.842078
17	18	0.842593
18	19	0.835905
19	20	0.834877
20	21	0.828189
21	22	0.827160
22	23	0.823560
23	24	0.819444
24	25	0.816358
25	26	0.821502
26	27	0.812243
27	28	0.815329
28	29	0.814300
29	30	0.814815
30	31	0.810700
31	32	0.812757
32	33	0.810105

```

32 33 0.810185
33 34 0.812757
34 35 0.810700
35 36 0.811728
36 37 0.809671
37 38 0.813786
38 39 0.806070

```

In [258]:

```

from sklearn.metrics import mean_squared_error
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import roc_curve, auc
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix

knn4 = KNeighborsClassifier()
knn4.set_params(n_neighbors = 3)
knn_model = knn4.fit(X_train, y_train)
y_pred_knn = knn4.predict(X_test)
y_pred_knn

# Get the confusion matrix
cf_matrix_knn = confusion_matrix(y_test, y_pred_knn)

print('Confusion Matrix - KNN')
print(cf_matrix_knn)

cf_matrix_plot = sns.heatmap(cf_matrix_knn, annot=True, cmap='Blues', annot_kws={"fontsize":18})

cf_matrix_plot.set_title('Confusion Matrix for KNN\n\n');
cf_matrix_plot.set_xlabel('\nStroke prediction')
cf_matrix_plot.set_ylabel('Actual Stroke condition ');

# Ticket labels
cf_matrix_plot.xaxis.set_ticklabels(['Stroke=No', 'Stroke=Yes'], fontsize = 18)
cf_matrix_plot.yaxis.set_ticklabels(['Stroke=No', 'Stroke=Yes'], fontsize = 18)

# Display the visualization of the Confusion Matrix.
plt.show()

print(classification_report(y_test, y_pred_knn))

class_probabilities = knn_model.predict_proba(X_test)
preds = class_probabilities[:, 1]

fpr, tpr, threshold = roc_curve(y_test, preds)
roc_auc = auc(fpr, tpr)

# Printing AUC
print(f"AUC for KNN classifier is: {roc_auc}")

# Plotting the ROC
plt.title('Receiver Operating Characteristic(KNN model)')
plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()

```

Confusion Matrix - KNN


```

[[1183  15]
 [  77   3]]

```


Confusion Matrix for KNN





	precision	recall	f1-score	support
0	0.94	0.99	0.96	1198
1	0.17	0.04	0.06	80
accuracy			0.93	1278
macro avg	0.55	0.51	0.51	1278
weighted avg	0.89	0.93	0.91	1278

AUC for KNN classifier is: 0.5874426126878131



In [259]:

```
#KNN model performance testing with oversampled data
knn4 = KNeighborsClassifier()
knn4.set_params(n_neighbors = 3)
knn_model = knn4.fit(X_train_os, y_train_os)
y_pred_knn = knn4.predict(X_test_os)
y_pred_knn

# Get the confusion matrix
cf_matrix_knn = confusion_matrix(y_test_os, y_pred_knn)

print('Confusion Matrix - KNN')
print(cf_matrix_knn)

cf_matrix_plot = sns.heatmap(cf_matrix_knn, annot=True, cmap='Blues', annot_kws={"fontsize":18})

cf_matrix_plot.set_title('Confusion Matrix for KNN\n\n');
cf_matrix_plot.set_xlabel('\nStroke prediction');
cf_matrix_plot.set_ylabel('Actual Stroke condition ');

# Ticket labels
cf_matrix_plot.xaxis.set_ticklabels(['Stroke=No','Stroke=Yes'],fontsize = 18)
cf_matrix_plot.yaxis.set_ticklabels(['Stroke=No','Stroke=Yes'],fontsize = 18)

# Display the visualization of the Confusion Matrix.
plt.show()
```

```

print(classification_report(y_test_os, y_pred_knn))

class_probabilities = knn_model.predict_proba(X_test_os)
preds = class_probabilities[:, 1]

fpr, tpr, threshold = roc_curve(y_test_os, preds)
roc_auc = auc(fpr, tpr)

# Printing AUC
print(f"AUC for KNN classifier is: {roc_auc}")

# Plotting the ROC
plt.title('Receiver Operating Characteristic(KNN model)')
plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()

```

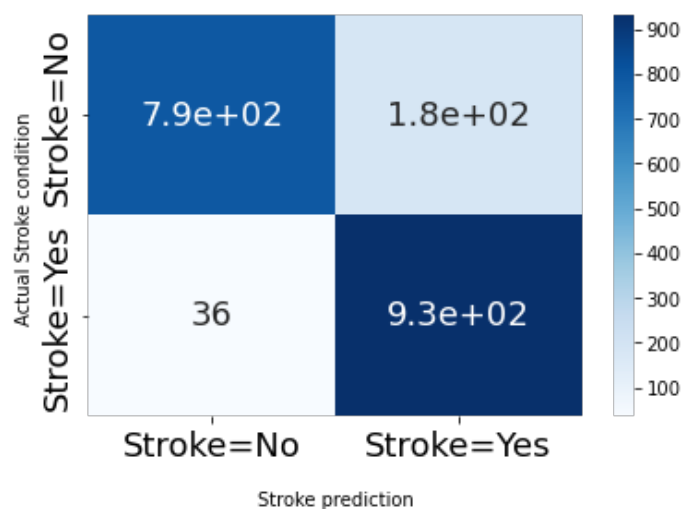
Confusion Matrix - KNN

```

[[793 183]
 [ 36 932]]


```


Confusion Matrix for KNN



	precision	recall	f1-score	support
0	0.96	0.81	0.88	976
1	0.84	0.96	0.89	968
accuracy			0.89	1944
macro avg	0.90	0.89	0.89	1944
weighted avg	0.90	0.89	0.89	1944

AUC for KNN classifier is: 0.9308375177821433





In [260]:

```
# Performing Hyper-parameter tuning to find the best model parameter. In case case, finding the best k value using GridSearchCV
from sklearn.model_selection import GridSearchCV
knn = KNeighborsClassifier()
params = {
    'n_neighbors': [3,5,7,9,11,13],
    'weights': ['uniform', 'distance'],
    'p': [1,2]
}

clf = GridSearchCV(
    estimator=knn,
    param_grid=params,
    cv=5,
    n_jobs=5,
    verbose=1
)
# Fitting our GridSearchCV Object
clf.fit(X_train, y_train)
# Printing the best parameters
print(clf.best_params_)
```

Fitting 5 folds for each of 24 candidates, totalling 120 fits  
{'n\_neighbors': 9, 'p': 2, 'weights': 'uniform'}

In [261]:

```
knn4 = KNeighborsClassifier()
knn4.set_params(n_neighbors = 5)
knn_model = knn4.fit(X_train, y_train)
y_pred_knn = knn4.predict(X_test)
y_pred_knn

# Get the confusion matrix
cf_matrix_knn = confusion_matrix(y_test, y_pred_knn)

print('Confusion Matrix - KNN')
print(cf_matrix_knn)

cf_matrix_plot = sns.heatmap(cf_matrix_knn, annot=True, cmap='Blues', annot_kws={"fontsize":18})

cf_matrix_plot.set_title('Confusion Matrix for KNN\n\n');
cf_matrix_plot.set_xlabel('\nStroke prediction')
cf_matrix_plot.set_ylabel('Actual Stroke condition ');

# Ticket labels
cf_matrix_plot.xaxis.set_ticklabels(['Stroke=No', 'Stroke=Yes'], fontsize = 18)
cf_matrix_plot.yaxis.set_ticklabels(['Stroke=No', 'Stroke=Yes'], fontsize = 18)

# Display the visualization of the Confusion Matrix.
plt.show()

print(classification_report(y_test, y_pred_knn))

class_probabilities = knn_model.predict_proba(X_test)
preds = class_probabilities[:, 1]

fpr, tpr, threshold = roc_curve(y_test, preds)
roc_auc = auc(fpr, tpr)

# Printing AUC
print(f"AUC for KNN classifier is: {roc_auc}")


# Plotting the ROC
plt.title('Receiver Operating Characteristic(KNN model)')
```

```
plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```

Confusion Matrix - KNN


```
[[1195    3]
 [  77    3]]
```

Confusion Matrix for KNN



	precision	recall	f1-score	support
0	0.94	1.00	0.97	1198
1	0.50	0.04	0.07	80
accuracy			0.94	1278
macro avg	0.72	0.52	0.52	1278
weighted avg	0.91	0.94	0.91	1278

AUC for KNN classifier is: 0.6586863522537563



In [262]:

```
# !pip install imbalanced-learn
```

In [263]:

```
# Logistic Regression Classifier
```

```
from sklearn.linear_model import LogisticRegression
lrc_classifier = LogisticRegression(random_state = 1)
```



```

lrc_model = lrc_classifier.fit(X_train, y_train)
y_pred_lrc = lrc_classifier.predict(X_test)
y_pred_lrc
# Get the confusion matrix
cf_matrix_lrc = confusion_matrix(y_test, y_pred_lrc)
print('Confusion Matrix - LRC')
print(cf_matrix_lrc)

cf_matrix_plot = sns.heatmap(cf_matrix_lrc, annot=True, cmap='Blues',annot_kws={"fontsize":18})

cf_matrix_plot.set_title('Confusion Matrix for Logistic Regression Classifier\n\n');
cf_matrix_plot.set_xlabel('\nStroke prediction')
cf_matrix_plot.set_ylabel('Actual Stroke condition ');

# Ticket labels
cf_matrix_plot.xaxis.set_ticklabels(['Stroke=No', 'Stroke=Yes'], fontsize = 18)
cf_matrix_plot.yaxis.set_ticklabels(['Stroke=No', 'Stroke=Yes'], fontsize = 18)

# Display the visualization of the Confusion Matrix.
plt.show()

print(classification_report(y_test, y_pred_lrc))

class_probabilities = lrc_model.predict_proba(X_test)
preds = class_probabilities[:, 1]

fpr, tpr, threshold = roc_curve(y_test, preds)
roc_auc = auc(fpr, tpr)

# Printing AUC
print(f"AUC for Logistic Regression classifier is: {roc_auc}")

# Plotting the ROC
plt.title('Receiver Operating Characteristic(LRC model)')
plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()


```

```

Confusion Matrix - LRC
[[1198    0]
 [  79    1]]


```

Confusion Matrix for Logistic Regression Classifier



accuracy			0.94	1278
macro avg	0.97	0.51	0.50	1278
weighted avg	0.94	0.94	0.91	1278

AUC for Logistic Regression classifier is: 0.8351106010016695



In [264]:

```
# LRC model performance testing with oversampled data

lrc_classifier = LogisticRegression(random_state = 1)
lrc_model = lrc_classifier.fit(X_train_os, y_train_os)
y_pred_lrc = lrc_classifier.predict(X_test_os)
y_pred_lrc
# Get the confusion matrix
cf_matrix_lrc = confusion_matrix(y_test_os, y_pred_lrc)
print('Confusion Matrix - LRC')
print(cf_matrix_lrc)

cf_matrix_plot = sns.heatmap(cf_matrix_lrc, annot=True, cmap='Blues', annot_kws={"fontsize":18})

cf_matrix_plot.set_title('Confusion Matrix for Logistic Regression Classifier\n\n');
cf_matrix_plot.set_xlabel('\nStroke prediction')
cf_matrix_plot.set_ylabel('Actual Stroke condition ');

# Ticket labels
cf_matrix_plot.xaxis.set_ticklabels(['Stroke=No', 'Stroke=Yes'], fontsize = 18)
cf_matrix_plot.yaxis.set_ticklabels(['Stroke=No', 'Stroke=Yes'], fontsize = 18)

# Display the visualization of the Confusion Matrix.
plt.show()

print(classification_report(y_test_os, y_pred_lrc))

class_probabilities = lrc_model.predict_proba(X_test_os)
preds = class_probabilities[:, 1]

fpr, tpr, threshold = roc_curve(y_test_os, preds)
roc_auc = auc(fpr, tpr)

# Printing AUC
print(f"AUC for Logistic Regression classifier is: {roc_auc}")


# Plotting the ROC
plt.title('Receiver Operating Characteristic(LRC model)')
plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
```

```
plt.show()
```

Confusion Matrix - LRC


```
[[726 250]
 [161 807]]
```

Confusion Matrix for Logistic Regression Classifier



	precision	recall	f1-score	support
0	0.82	0.74	0.78	976
1	0.76	0.83	0.80	968
accuracy			0.79	1944
macro avg	0.79	0.79	0.79	1944
weighted avg	0.79	0.79	0.79	1944

AUC for Logistic Regression classifier is: 0.8498763717653435



In [265]:

```
# SVM Classifier

from sklearn.svm import SVC
classifier_svc = SVC(kernel = 'linear', random_state = 0,probability=True) #class_weight
='balanced',
svm_model = classifier_svc.fit(X_train, y_train)
y_pred_svm = classifier_svc.predict(X_test)
y_pred_svm
# Get the confusion matrix
cf_matrix_svm = confusion_matrix(y_test, y_pred_svm)
print('Confusion Matrix - SVM')
print(cf_matrix_svm)

cf_matrix_plot = sns.heatmap(cf_matrix_svm, annot=True, cmap='Blues',annot_kws={"fontsize
":18})
```

```

cf_matrix_plot.set_title('Confusion Matrix for SVM Classifier\n\n');
cf_matrix_plot.set_xlabel('\nStroke prediction')
cf_matrix_plot.set_ylabel('Actual Stroke condition ');

# Ticket labels
cf_matrix_plot.xaxis.set_ticklabels(['Stroke=No', 'Stroke=Yes'], fontsize = 18)
cf_matrix_plot.yaxis.set_ticklabels(['Stroke=No', 'Stroke=Yes'], fontsize = 18)

# Display the visualization of the Confusion Matrix.
plt.show()

print(classification_report(y_test, y_pred_svm))

class_probabilities = svm_model.predict_proba(X_test)
preds = class_probabilities[:, 1]

fpr, tpr, threshold = roc_curve(y_test, preds)
roc_auc = auc(fpr, tpr)

# Printing AUC
print(f"AUC for SVM classifier is: {roc_auc}")


# Plotting the ROC
plt.title('Receiver Operating Characteristic(SVM model)')
plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()

```

```

Confusion Matrix - SVM
[[1198    0]
 [  80    0]]

```



```

AUC for SVM classifier is: 0.4173935726210351

```

```

C:\Users\samru\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning:

```


Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

C:\Users\samru\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318: UndefinedMetricWarning:

Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

C:\Users\samru\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318: UndefinedMetricWarning:

Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.



In [266]:

```
# SVM model performance testing with oversampled data

classifier_svc = SVC(kernel = 'linear', random_state = 0, probability=True) #class_weight
='balanced',
svm_model = classifier_svc.fit(X_train_os, y_train_os)
y_pred_svm = classifier_svc.predict(X_test_os)
y_pred_svm

# Get the confusion matrix
cf_matrix_svm = confusion_matrix(y_test_os, y_pred_svm)
print('Confusion Matrix - SVM')
print(cf_matrix_svm)

cf_matrix_plot = sns.heatmap(cf_matrix_svm, annot=True, cmap='Blues', annot_kws={"fontsize
":18})

cf_matrix_plot.set_title('Confusion Matrix for SVM Classifier\n\n');
cf_matrix_plot.set_xlabel('\nStroke prediction')
cf_matrix_plot.set_ylabel('Actual Stroke condition ');

# Ticket labels
cf_matrix_plot.xaxis.set_ticklabels(['Stroke=No', 'Stroke=Yes'], fontsize = 18)
cf_matrix_plot.yaxis.set_ticklabels(['Stroke=No', 'Stroke=Yes'], fontsize = 18)

# Display the visualization of the Confusion Matrix.
plt.show()

print(classification_report(y_test_os, y_pred_svm))

class_probabilities = svm_model.predict_proba(X_test_os)
preds = class_probabilities[:, 1]

fpr, tpr, threshold = roc_curve(y_test_os, preds)
roc_auc = auc(fpr, tpr)

# Printing AUC
```


```
print(f"AUC for SVM classifier is: {roc_auc}")
```

```
# Plotting the ROC
plt.title('Receiver Operating Characteristic(SVM model)')
plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```

Confusion Matrix - SVM


```
[[711 265]
 [127 841]]
```

Confusion Matrix for SVM Classifier



	precision	recall	f1-score	support
0	0.85	0.73	0.78	976
1	0.76	0.87	0.81	968
accuracy			0.80	1944
macro avg	0.80	0.80	0.80	1944
weighted avg	0.80	0.80	0.80	1944

AUC for SVM classifier is: 0.8517853060222191



In [267]:

```
# !pip install mlxtend
```

In [268]:

```
# ...
```

```

# Random Forest Classifier

from sklearn.ensemble import RandomForestClassifier
classifier_rf = RandomForestClassifier( random_state=0)
rf_model = classifier_rf.fit(X_train, y_train)
y_pred_rf = classifier_rf.predict(X_test)
y_pred_rf
# Get the confusion matrix
cf_matrix_rf = confusion_matrix(y_test, y_pred_rf)
print('Confusion Matrix - RF')
print(cf_matrix_rf)

# Plotting Confusion matrix
cf_matrix_plot = sns.heatmap(cf_matrix_rf, annot=True, cmap='Blues',annot_kws={"fontsize":18})
cf_matrix_plot.set_title('Confusion Matrix for Random Forest\n\n');
cf_matrix_plot.set_xlabel('\nStroke prediction')
cf_matrix_plot.set_ylabel('Actual Stroke condition ');

# Ticket labels - List must be in alphabetical order
cf_matrix_plot.xaxis.set_ticklabels(['Stroke=No', 'Stroke=Yes'], fontsize = 18)
cf_matrix_plot.yaxis.set_ticklabels(['Stroke=No', 'Stroke=Yes'], fontsize = 18)
plt.show()

print(classification_report(y_test, y_pred_rf))

class_probabilities = rf_model.predict_proba(X_test)
preds = class_probabilities[:, 1]

fpr, tpr, threshold = roc_curve(y_test, preds)
roc_auc = auc(fpr, tpr)

# Printing AUC
print(f"AUC for Random Forest classifier is: {roc_auc}")


# Plotting the ROC
plt.title('Receiver Operating Characteristic(Random Forest model)')
plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()

```

Confusion Matrix - RF

```
[[1197    1]
 [   80    0]]
```


Confusion Matrix for Random Forest



	precision	recall	f1-score	support
Stroke=No	0.99	1.00	0.99	1198
Stroke=Yes	0.00	0.00	0.00	80

	0	0.94	1.00	0.97	1198
	1	0.00	0.00	0.00	80
accuracy				0.94	1278
macro avg		0.47	0.50	0.48	1278
weighted avg		0.88	0.94	0.91	1278

AUC for Random Forest classifier is: 0.7967810934891485



In [269]:

```
# Random Forest Classifier performance testing after oversampling

classifier_rf = RandomForestClassifier( random_state=0)
rf_model = classifier_rf.fit(X_train_os, y_train_os)
y_pred_rf = classifier_rf.predict(X_test_os)
y_pred_rf
# Get the confusion matrix
cf_matrix_rf = confusion_matrix(y_test_os, y_pred_rf)
print('Confusion Matrix - RF')
print(cf_matrix_rf)

# Plotting Confusion matrix
cf_matrix_plot = sns.heatmap(cf_matrix_rf, annot=True, cmap='Blues',annot_kws={"fontsize":18})
cf_matrix_plot.set_title('Confusion Matrix for Random Forest\n\n');
cf_matrix_plot.set_xlabel('\nStroke prediction')
cf_matrix_plot.set_ylabel('Actual Stroke condition ');

# Ticket labels - List must be in alphabetical order
cf_matrix_plot.xaxis.set_ticklabels(['Stroke=No','Stroke=Yes'],fontsize = 18)
cf_matrix_plot.yaxis.set_ticklabels(['Stroke=No','Stroke=Yes'],fontsize = 18)
plt.show()

print(classification_report(y_test_os, y_pred_rf))

class_probabilities = rf_model.predict_proba(X_test_os)
preds = class_probabilities[:, 1]

fpr, tpr, threshold = roc_curve(y_test_os, preds)
roc_auc = auc(fpr, tpr)

# Printing AUC
print(f"AUC for Random Forest classifier is: {roc_auc}")

# Plotting the ROC
plt.title('Receiver Operating Characteristic(Random forest model)')
plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
```




```
plt.show()
```

Confusion Matrix - RF


```
[[868 108]
 [ 26 942]]
```

Confusion Matrix for Random Forest



	precision	recall	f1-score	support
0	0.97	0.89	0.93	976
1	0.90	0.97	0.93	968
accuracy			0.93	1944
macro avg	0.93	0.93	0.93	1944
weighted avg	0.93	0.93	0.93	1944

AUC for Random Forest classifier is: 0.982222090502642



In [270]:

```
# Decision Tree Classifier

from sklearn.tree import DecisionTreeClassifier
from sklearn import tree

dt = DecisionTreeClassifier(random_state= 0)
dt_model = dt.fit(X_train, y_train)
y_pred_dt = dt.predict(X_test)
y_pred_dt

# Get the confusion matrix
cf_matrix_dt = confusion_matrix(y_test, y_pred_dt)
print('Confusion Matrix - RF')
print(cf_matrix_dt)
```

```

cf_matrix_plot = sns.heatmap(cf_matrix_dt, annot=True, cmap='Blues',annot_kws={"fontsize":18})

cf_matrix_plot.set_title('Confusion Matrix for Decision Tree\n\n');
cf_matrix_plot.set_xlabel('\nStroke prediction')
cf_matrix_plot.set_ylabel('Actual Stroke condition ');

## Ticket labels - List must be in alphabetical order
cf_matrix_plot.xaxis.set_ticklabels(['Stroke=No', 'Stroke=Yes'],fontsize = 18)
cf_matrix_plot.yaxis.set_ticklabels(['Stroke=No', 'Stroke=Yes'],fontsize = 18)

## Display the visualization of the Confusion Matrix.
plt.show()

print(classification_report(y_test, y_pred_dt))

class_probabilities = dt_model.predict_proba(X_test)
preds = class_probabilities[:, 1]

fpr, tpr, threshold = roc_curve(y_test, preds)
roc_auc = auc(fpr, tpr)

# Printing AUC
print(f"AUC for Decision Trees classifier is: {roc_auc}")

# Plotting the ROC
plt.title('Receiver Operating Characteristic(Decision Trees model)')
plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
# plt.figure(figsize=(50,50))
# tree.plot_tree(dt, filled= True,feature_names=X_train.columns, fontsize=10)

```


Confusion Matrix - RF

```

[[1151   47]
 [  69   11]]


```

Confusion Matrix for Decision Tree



	precision	recall	f1-score	support
0	0.94	0.96	0.95	1198
1	0.19	0.14	0.16	80
accuracy			0.91	1278
macro avg	0.57	0.55	0.56	1278
weighted avg	0.90	0.91	0.90	1278

AUC for Decision Trees classifier is: 0.5491339732888147



In [271]:

```
# Decision Trees Classifier performance testing after oversampling
dt = DecisionTreeClassifier(random_state= 0)
dt_model = dt.fit(X_train_os, y_train_os)
y_pred_dt = dt.predict(X_test_os)
y_pred_dt

# Get the confusion matrix
cf_matrix_dt = confusion_matrix(y_test_os, y_pred_dt)
print('Confusion Matrix - RF')
print(cf_matrix_dt)

cf_matrix_plot = sns.heatmap(cf_matrix_dt, annot=True, cmap='Blues',annot_kws={"fontsize":18})

cf_matrix_plot.set_title('Confusion Matrix for Decision Tree\n\n');
cf_matrix_plot.set_xlabel('\nStroke prediction')
cf_matrix_plot.set_ylabel('Actual Stroke condition ');

# Ticket labels - List must be in alphabetical order
cf_matrix_plot.xaxis.set_ticklabels(['Stroke=No', 'Stroke=Yes'], fontsize = 18)
cf_matrix_plot.yaxis.set_ticklabels(['Stroke=No', 'Stroke=Yes'], fontsize = 18)

# Display the visualization of the Confusion Matrix.
plt.show()

print(classification_report(y_test_os, y_pred_dt))

class_probabilities = dt_model.predict_proba(X_test_os)
preds = class_probabilities[:, 1]


fpr, tpr, threshold = roc_curve(y_test_os, preds)
roc_auc = auc(fpr, tpr)

# Printing AUC
print(f"AUC for Decision Trees classifier is: {roc_auc}")

# Plotting the ROC
plt.title('Receiver Operating Characteristic(Decision Trees model)')
plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```


```
Confusion Matrix - RF
[[852 124]
 [115 853]]
```

Confusion Matrix for Decision Tree



	precision	recall	f1-score	support
0	0.88	0.87	0.88	976
1	0.87	0.88	0.88	968
accuracy			0.88	1944
macro avg	0.88	0.88	0.88	1944
weighted avg	0.88	0.88	0.88	1944

AUC for Decision Trees classifier is: 0.8770745833897846



In [272]:

```
from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import chi2
# selecting top 10 features using chi-square test method
sel_chi2 = SelectKBest(chi2, k=10)
X_train_chi2 = sel_chi2.fit_transform(X_train, y_train)
print(sel_chi2.get_support())
```

```
[False  True  True  True  True  True  True  False  True  False  True  True
  True  False  False]
```

In [273]:

```
X_test_chi2 = sel_chi2.transform(X_test)
print(X_test.shape)
print(X_test_chi2.shape)
```

```
(1278, 15)
(1278, 10)
```

In [274]:

```
# Random Forest Classifier performance testing after feature selection
```

```
classifier_rf = RandomForestClassifier( random_state=0)
rf_model = classifier_rf.fit(X_train_chi2, y_train)
y_pred_rf = classifier_rf.predict(X_test_chi2)
y_pred_rf
# Get the confusion matrix
cf_matrix_rf = confusion_matrix(y_test, y_pred_rf)
print('Confusion Matrix - RF')
print(cf_matrix_rf)

# Plotting Confusion matrix
cf_matrix_plot = sns.heatmap(cf_matrix_rf, annot=True, cmap='Blues',annot_kws={"fontsize":18})
cf_matrix_plot.set_title('Confusion Matrix for Random Forest\n\n');
cf_matrix_plot.set_xlabel('\nStroke prediction')
cf_matrix_plot.set_ylabel('Actual Stroke condition ');

# Ticket labels - List must be in alphabetical order
cf_matrix_plot.xaxis.set_ticklabels(['Stroke=No', 'Stroke=Yes'],fontsize = 18)
cf_matrix_plot.yaxis.set_ticklabels(['Stroke=No', 'Stroke=Yes'],fontsize = 18)
plt.show()

print(classification_report(y_test, y_pred_rf))

class_probabilities = rf_model.predict_proba(X_test_chi2)
preds = class_probabilities[:, 1]

fpr, tpr, threshold = roc_curve(y_test, preds)
roc_auc = auc(fpr, tpr)


# Printing AUC
print(f"AUC for Random Forest classifier is: {roc_auc}")

# Plotting the ROC
plt.title('Receiver Operating Characteristic(Random forest model)')
plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```

Confusion Matrix - RF


[[1186	12]
[ 77	3]]

Confusion Matrix for Random Forest



accuracy			0.93	1278
macro avg	0.57	0.51	0.51	1278
weighted avg	0.89	0.93	0.91	1278

AUC for Random Forest classifier is: 0.7708785475792989



In [ ]: